

10707

Deep Learning

Russ Salakhutdinov

Machine Learning Department

rsalakhu@cs.cmu.edu

Language Modeling

Natural Language Processing

- Natural language processing is concerned with tasks involving language data
 - we will focus on text data NLP
- Much like for computer vision, we can design neural networks specifically adapted to the processing of text data
 - main issue: text data is inherently **high dimensional**

Natural Language Processing

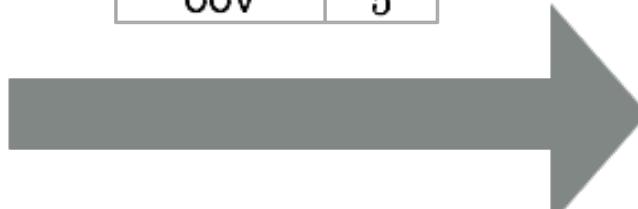
- Typical preprocessing steps of text data
 - Form vocabulary of words that maps words to a unique ID
 - Different criteria can be used to select which words are part of the vocabulary
 - Pick **most frequent words** and ignore **uninformative** words from a user-defined short list (ex.: “the”, “a”, etc.)
 - All words not in the vocabulary will be mapped to a special “**out-of-vocabulary**”
- Typical vocabulary sizes will vary between 100,000 and 1,000,000

Vocabulary

- Example:

" the "
" cat "
" and "
" the "
" dog "
" play "
" • "

Word	<i>w</i>
" the "	1
" and "	2
" dog "	3
" • "	4
" oov "	5



1
5
2
1
3
5
4

- We will note word IDs with the symbol w

- we can think of w as a **categorical feature** for the original word
- we will sometimes refer to w as a word, for simplicity

One-Hot Encoding

- From its word ID, we get a basic representation of a word through the **one-hot encoding** of the ID
 - the **one-hot vector** of an ID is a vector filled with 0s, except for a 1 at the position associated with the ID
 - For vocabulary size D=10, the one-hot vector of word ID w=4 is:
$$e(w) = [0\ 0\ 0\ 1\ 0\ 0\ 0\ 0\ 0\ 0]$$
- A one-hot encoding makes no assumption about **word similarity**
- This is a natural representation to start with, though a poor one

One-Hot Encoding

- The major problem with the one-hot representation is that it is very high-dimensional
 - the dimensionality of $e(w)$ is the size of the vocabulary
 - a typical vocabulary size is $\approx 100,000$
 - a window of 10 words would correspond to an input vector of **at least 1,000,000 units!**
- This has 2 consequences:
 - vulnerability to **overfitting** (millions of inputs means millions of parameters to train)
 - computationally **expensive**

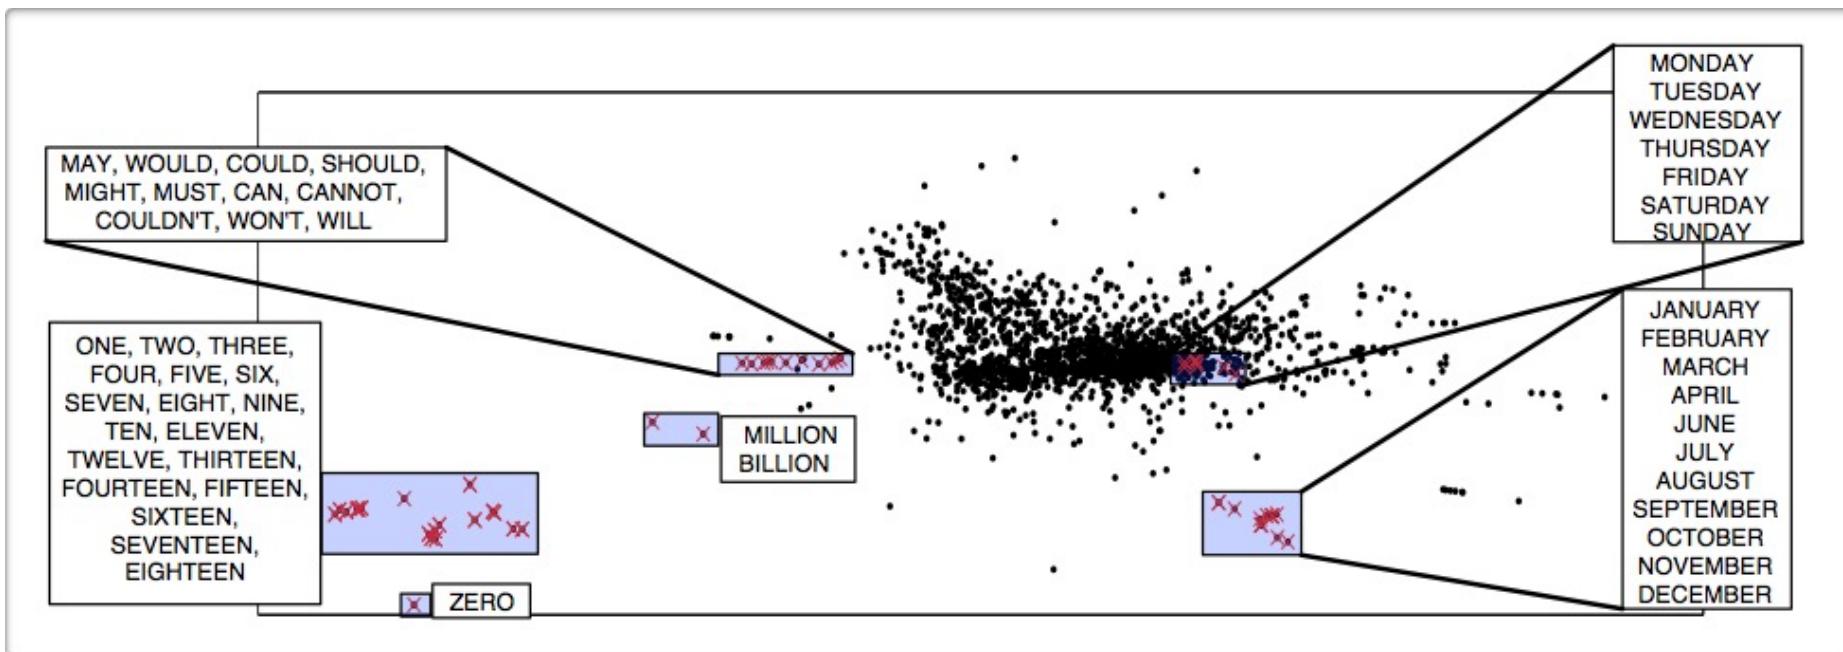
Continuous Representation of Words

- Each word w is associated with a real-valued vector $C(w)$

Word	w	$C(w)$
" the "	1	[0.6762, -0.9607, 0.3626, -0.2410, 0.6636]
" a "	2	[0.6859, -0.9266, 0.3777, -0.2140, 0.6711]
" have "	3	[0.1656, -0.1530, 0.0310, -0.3321, -0.1342]
" be "	4	[0.1760, -0.1340, 0.0702, -0.2981, -0.1111]
" cat "	5	[0.5896, 0.9137, 0.0452, 0.7603, -0.6541]
" dog "	6	[0.5965, 0.9143, 0.0899, 0.7702, -0.6392]
" car "	7	[-0.0069, 0.7995, 0.6433, 0.2898, 0.6359]
...

Continuous Representation of Words

- We would like the distance $\|C(w) - C(w')\|$ to reflect **meaningful similarities** between words



(from Blitzer et al. 2004)

Continuous Representation of Words

- Learn a continuous representation of words
 - we could then use these representations as input to a neural network
- We learn these representations by **gradient descent**
 - we don't only update the neural network parameters
 - we also update **each representation $C(w)$** in the input x with a gradient step:

$$C(w) \leftarrow C(w) - \alpha \nabla_{C(w)} l$$

where l is the **loss function** optimized by the neural network

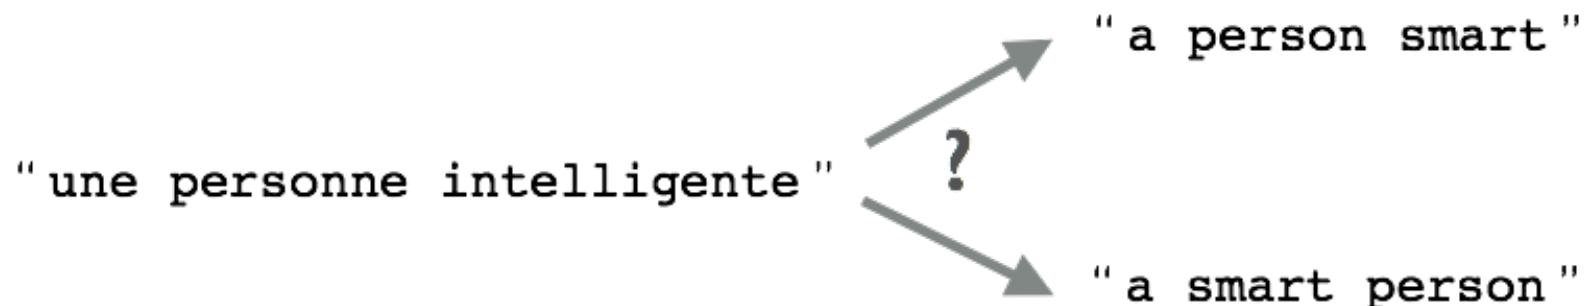
Continuous Representation of Words

- Let C be a matrix whose rows are the representations $C(w)$
 - obtaining $C(w)$ corresponds to the multiplication $e(w)^T C$
 - view differently, we are **projecting** $e(w)$ onto the columns of C
 - this is a **continuous transformation**, through which we can propagate gradients
- In practice, we implement $C(w)$ with a lookup table, not with a multiplication

Language Modeling

$$p(w_1, \dots, w_T)$$

- language modeling is the task of learning a language model that assigns **high probabilities** to well formed sentences
- plays a crucial role in speech recognition and machine translation systems



Language Modeling

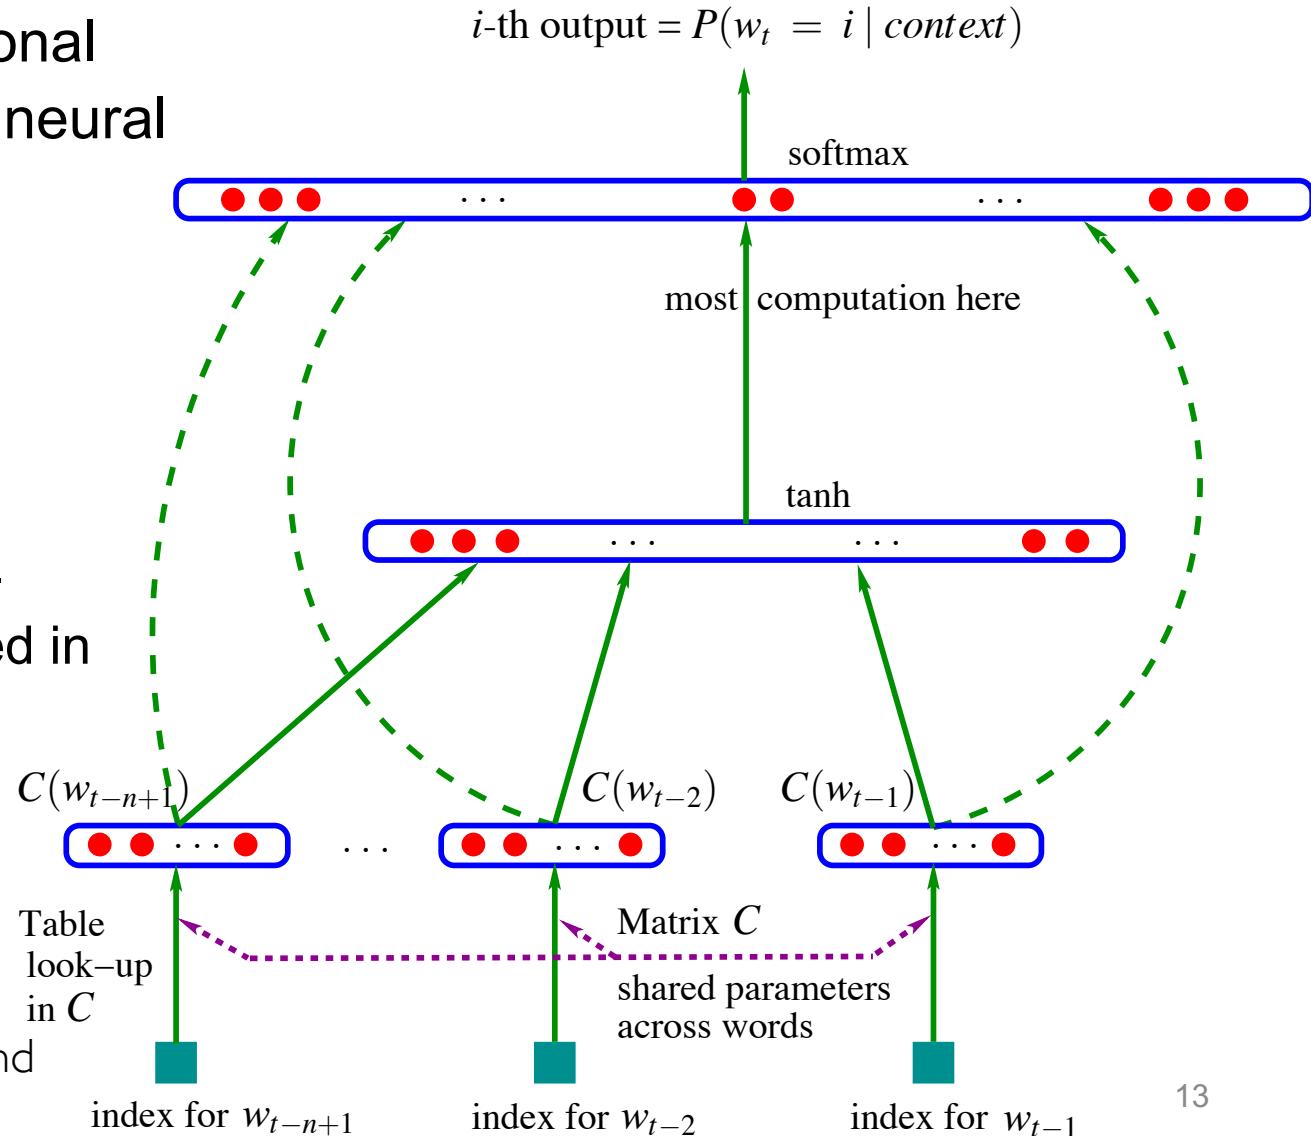
$$p(w_1, \dots, w_T) = \prod_{t=1}^T p(w_t | w_{t-(n-1)}, \dots, w_{t-1})$$

- the t^{th} word was generated based only on the $n-1$ previous words
- we will refer to $w_{t-(n-1)}, \dots, w_{t-1}$ as the context

Neural Language Model

- Model the conditional distributions with a neural network:

- learn word representations to allow transfer to n-grams not observed in training corpus



Neural Language Model

- Can potentially **generalize** to contexts not seen in training set
 - Example: $P(\text{"eating"} | \text{"the"}, \text{"cat"}, \text{"is"})$
 - Imagine 4-gram [**“ the ”**, **“ cat ”**, **“ is ”**, **“ eating ”**] is not in training corpus, but [**“ the ”**, **“ dog ”**, **“ is ”**, **“ eating ”**] is
 - If the word representations of **“ cat ”** and **“ dog ”** are similar, then the neural network will be able to generalize to the case of **“ cat ”**

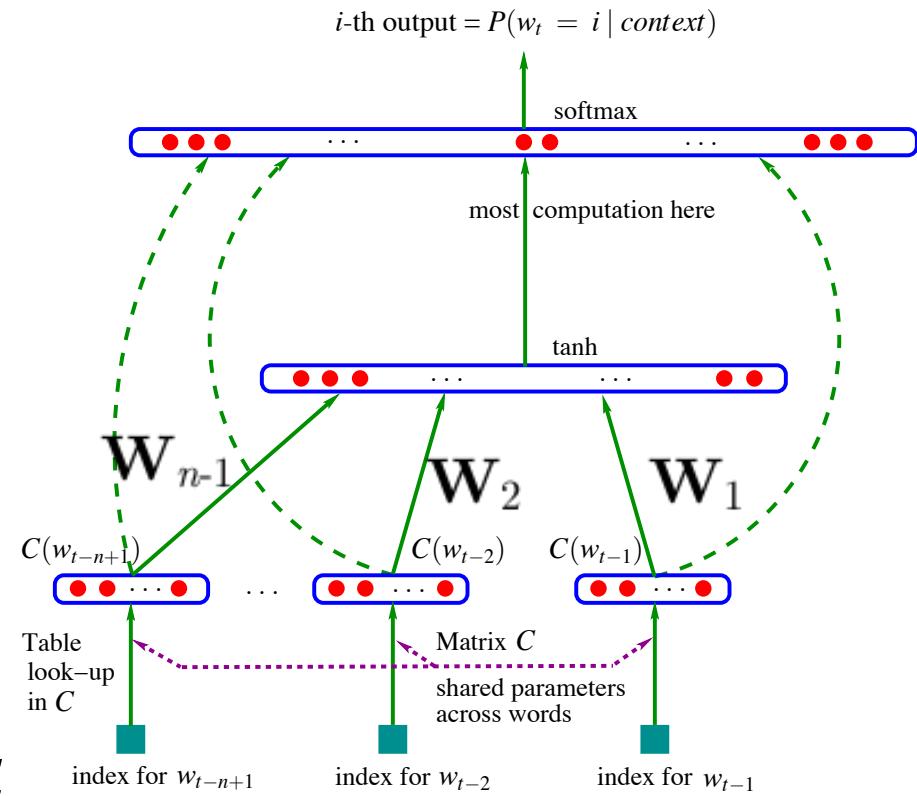
Neural Language Model

- We know how to propagate gradients in such a network

$$\nabla_{\mathbf{a}(\mathbf{x})} l$$

- let's note the submatrix connecting w_{t-i} and the hidden layer as \mathbf{W}_i
- The gradient wrt $C(w)$ for any w is

$$\nabla_{C(w)} l = \sum_{i=1}^{n-1} \mathbf{1}_{(w_{t-i}=w)} \mathbf{W}_i^\top \nabla_{\mathbf{a}(\mathbf{x})} l$$



Performance Evaluation

- In language modeling, a common evaluation metric is the **perplexity**
 - it is simply the exponential of the average negative log-likelihood
- Evaluation on Brown Corpus
 - n-gram model (Kneser-Ney smoothing): 321
 - neural network language model: 276
 - neural network + n-gram: 252

How About Generating Sentences!

Input



Output

A man skiing down the snow covered mountain with a dark sky in the background.

How About Generating Sentences!

Input



Output

A man skiing down the snow covered mountain with a dark sky in the background.

We want to model:

$$p(w_1, w_2, \dots, w_n) =$$

$$p(w_1)p(w_2|w_1)p(w_3|w_1, w_2)\dots p(w_n|w_1, w_2, \dots, w_{n-1})$$

Caption Generation with NLM



a car is parked in
the middle of nowhere .



a wooden table and chairs
arranged in a room .



a ferry boat on a marina
with a group of people .



there is a cat sitting on a shelf .



a little boy with a bunch
of friends on the street .

Caption Generation with NLM



the two birds are trying
to be seen in the water .
(can't count)



a giraffe is standing next
to a fence in a field .
(hallucination)



a parked car while
driving down the road .
(contradiction)

Caption Generation with NLM



the two birds are trying
to be seen in the water .
(can't count)



a giraffe is standing next
to a fence in a field .
(hallucination)



a parked car while
driving down the road .
(contradiction)



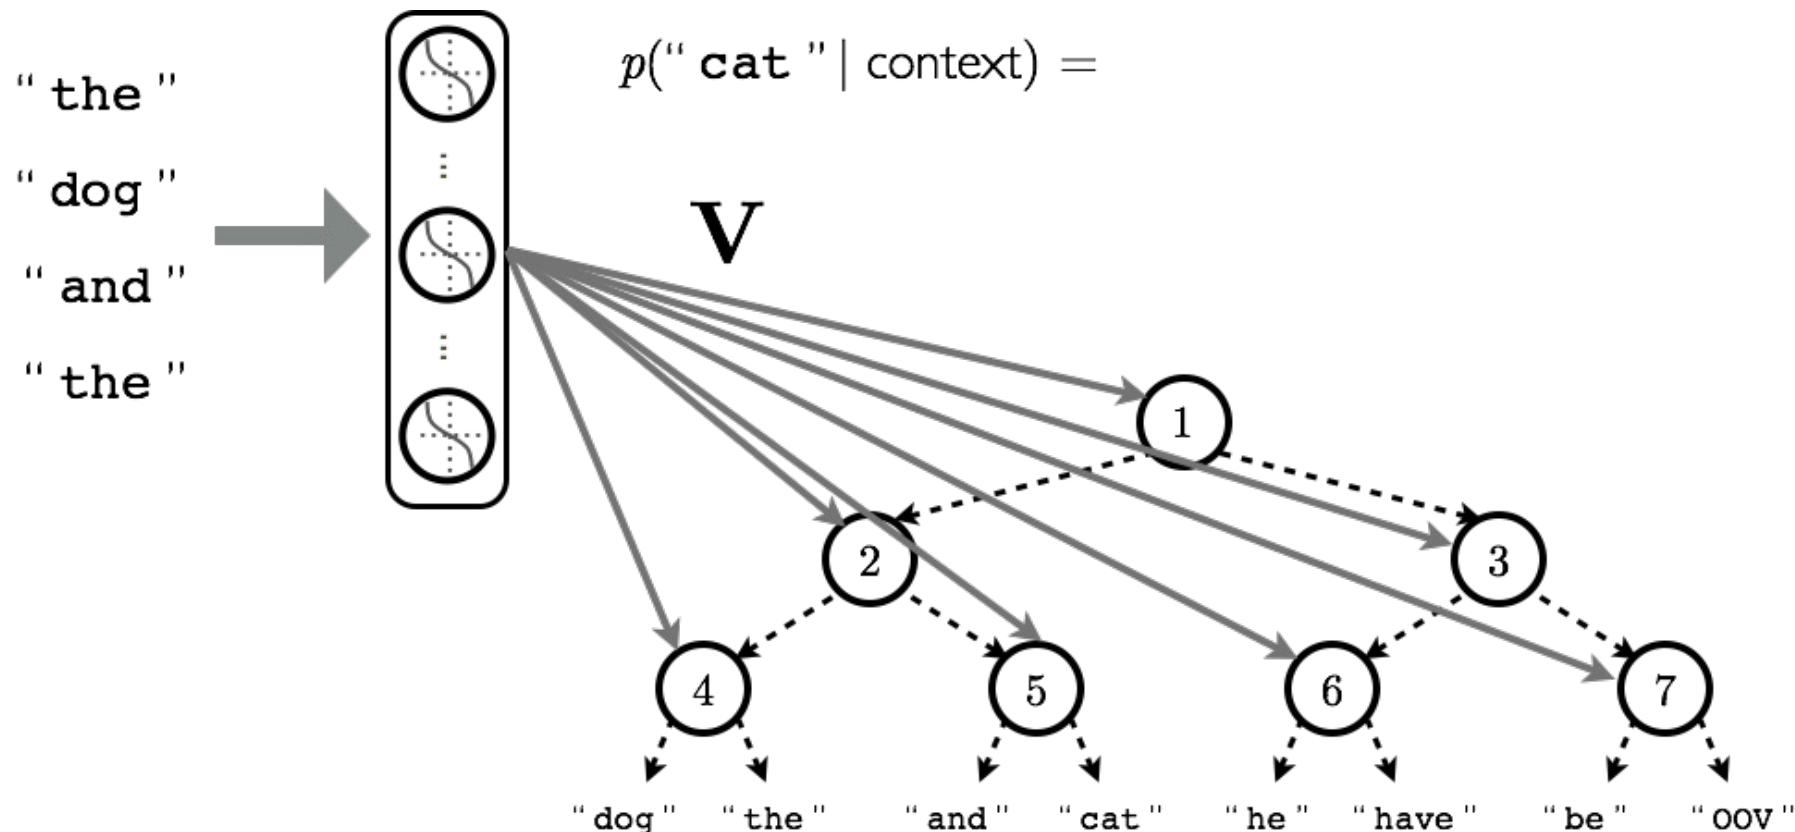
the handlebars are trying
to ride a bike rack .
(nonsensical)



a woman and a bottle of wine
in a garden .
(gender)

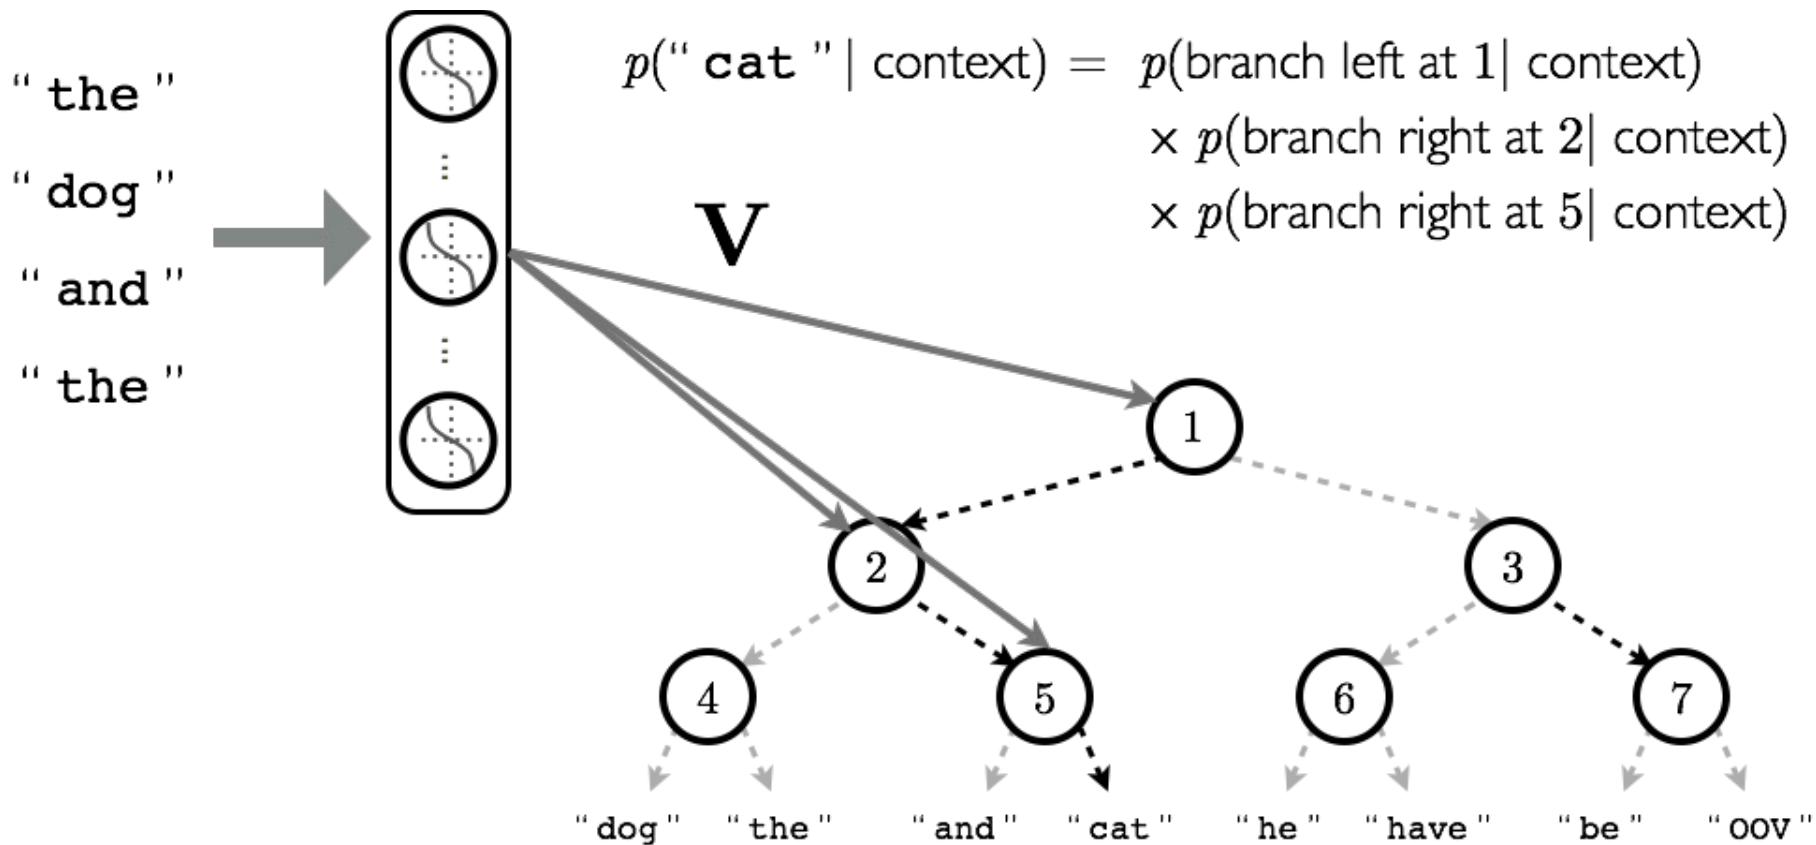
Hierarchical Output Layer

- Example: [“ the ”, “ dog ”, “ and ”, “ the ”, “ cat ”]



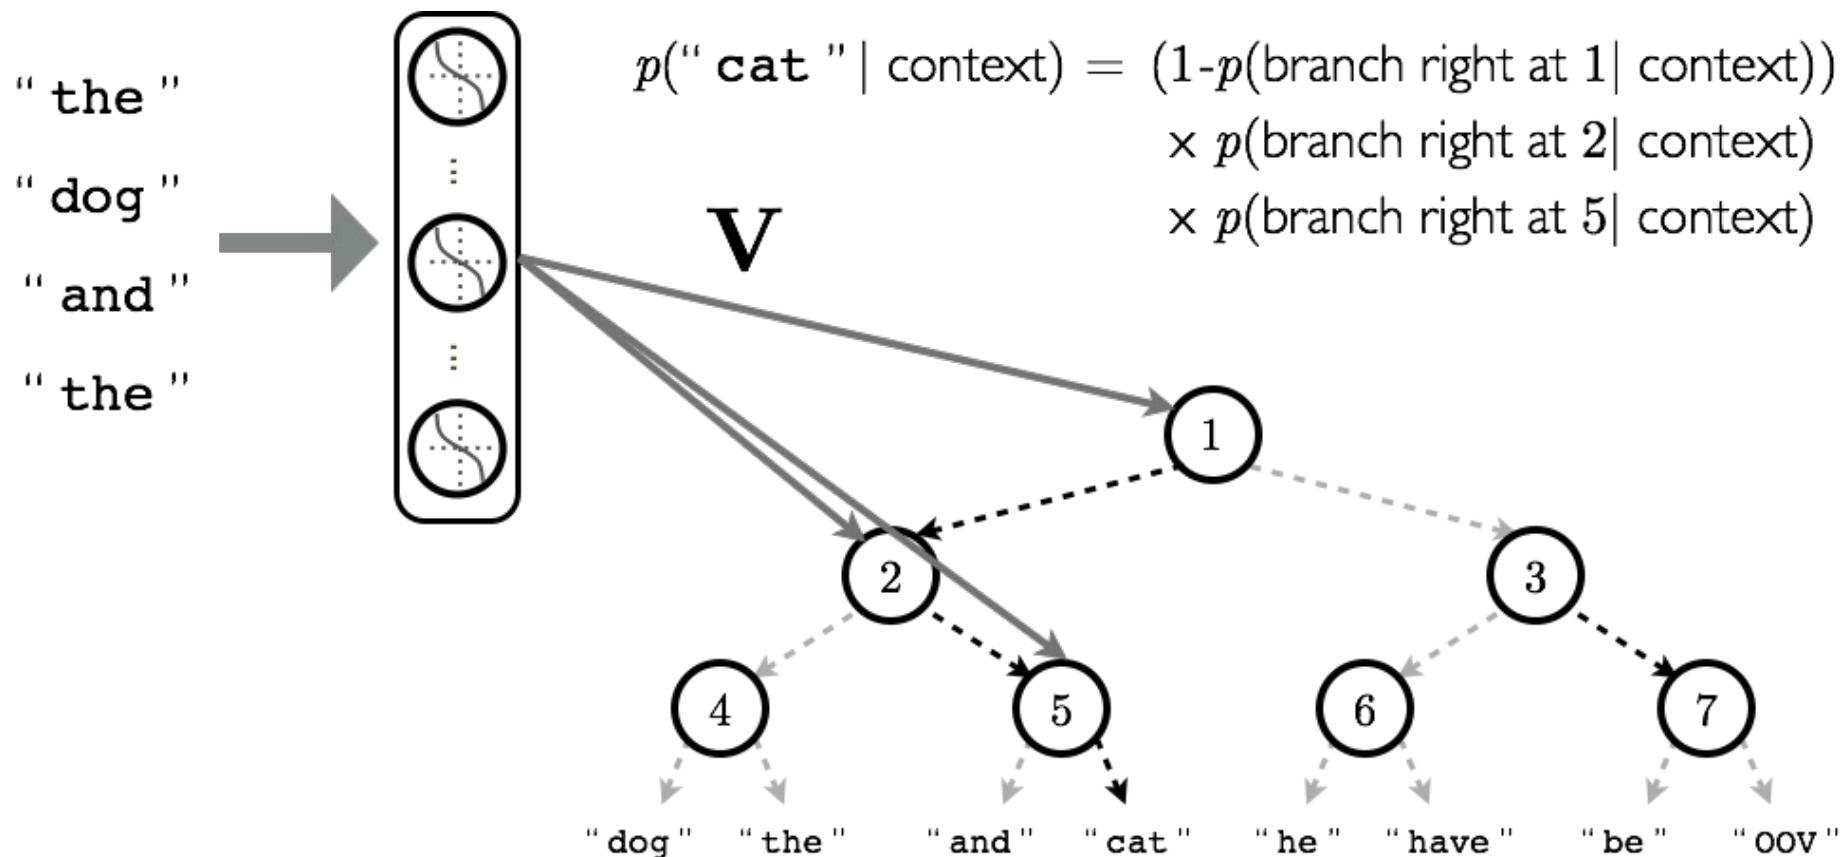
Hierarchical Output Layer

- Example: [“ the ”, “ dog ”, “ and ”, “ the ”, “ cat ”]



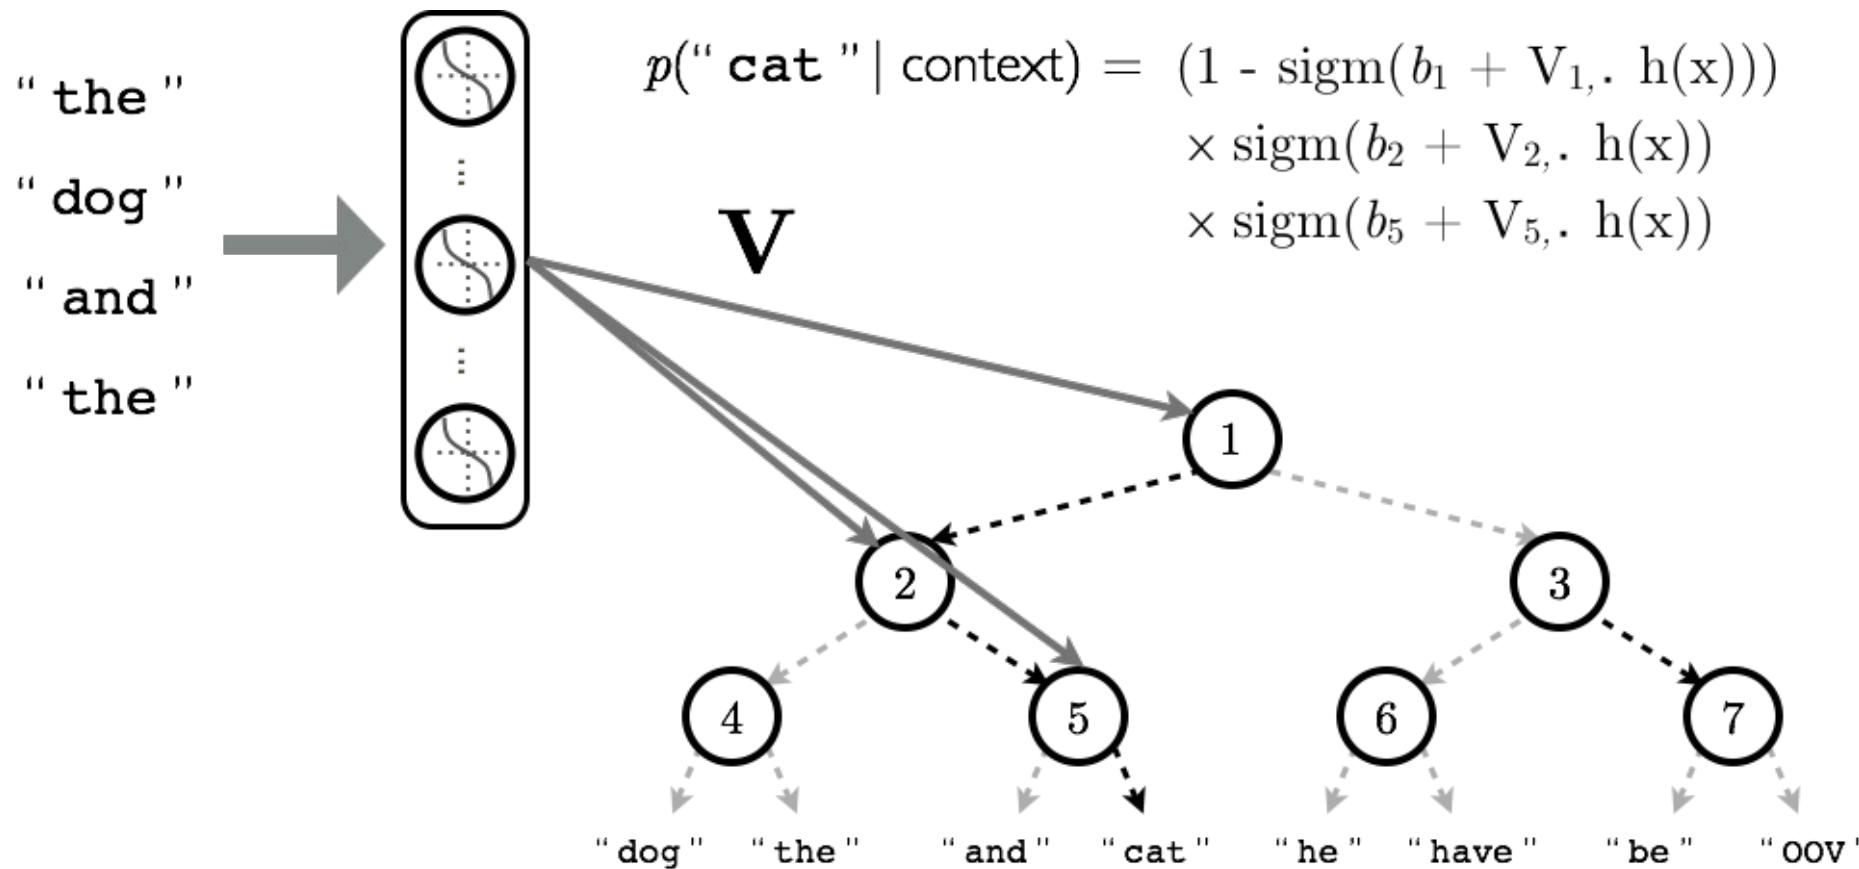
Hierarchical Output Layer

- Example: ["the", "dog", "and", "the", "cat"]



Hierarchical Output Layer

- Example: [“ the ”, “ dog ”, “ and ”, “ the ”, “ cat ”]



Hierarchical Output Layer

- How to define the word hierarchy?

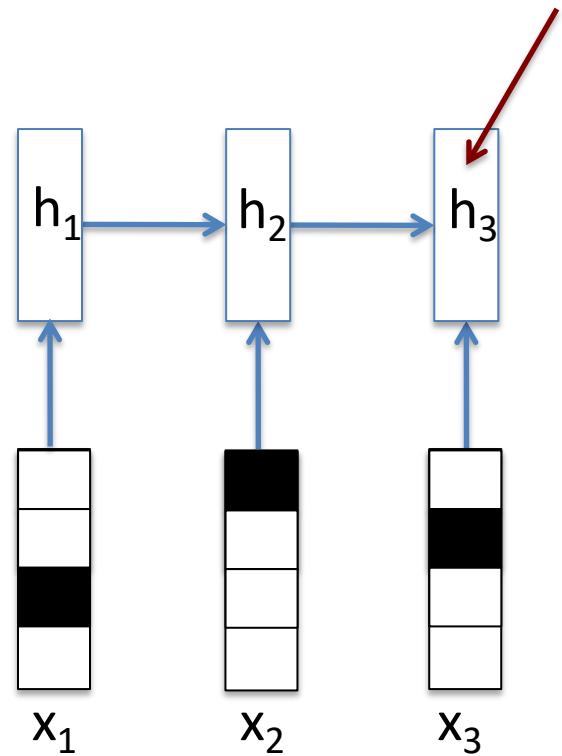
- can use a randomly generated tree
- can use existing linguistic resources, such as WordNet
- can learn the hierarchy using a recursive partitioning strategy

A Scalable Hierarchical Distributed Language Model Mnih and Hinton, 2008

They report a speedup of 100x, without performance decrease

Encoding Sentences via Recurrent Neural Network

Sentence
Representation

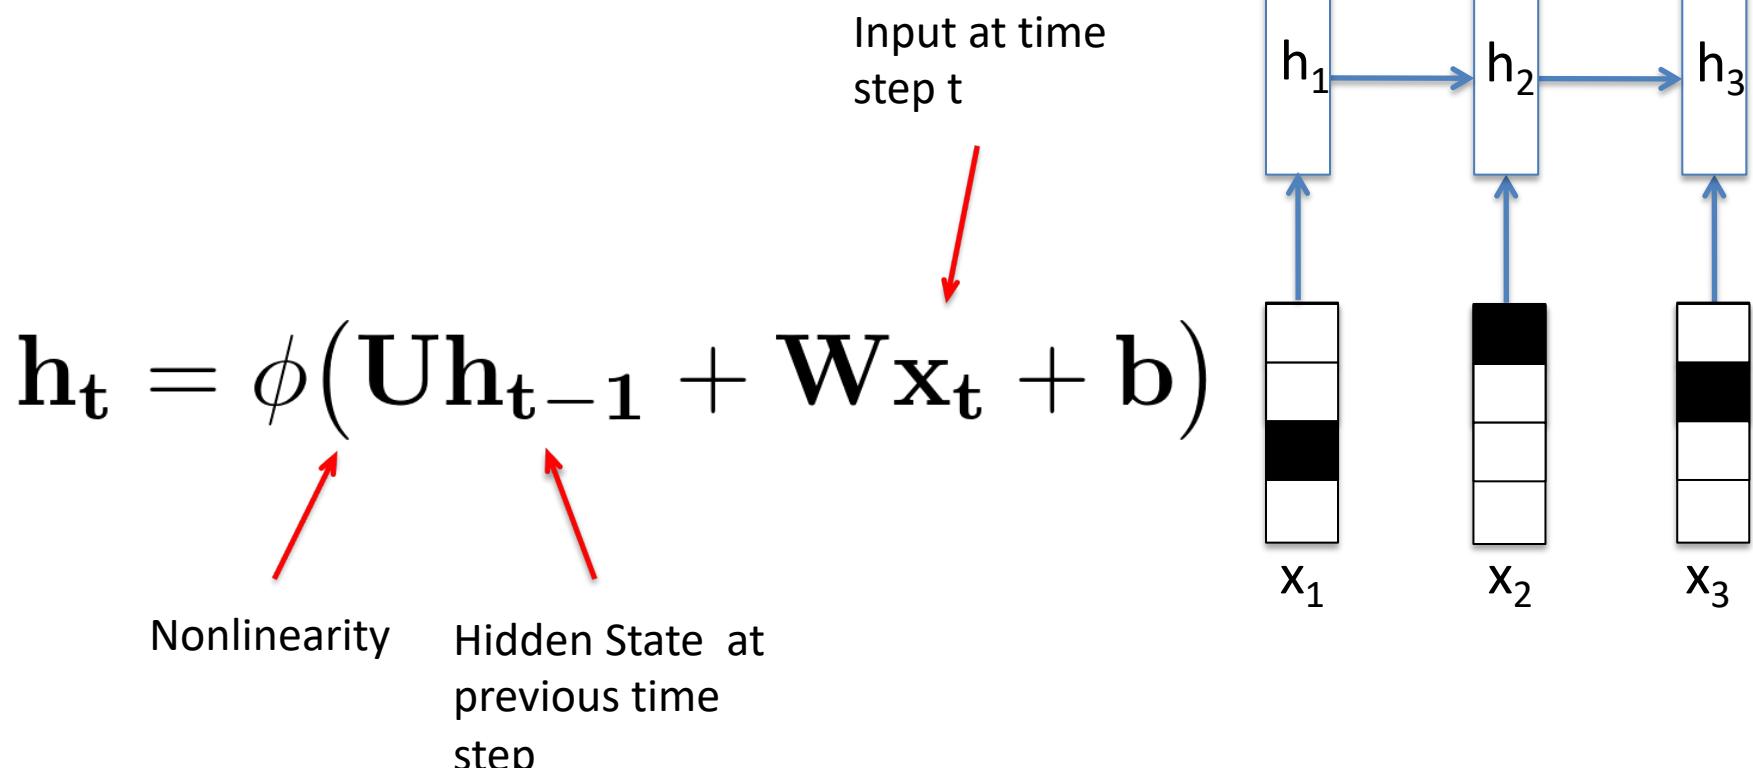


1-of-K encoding of words

Recurrent Neural Network

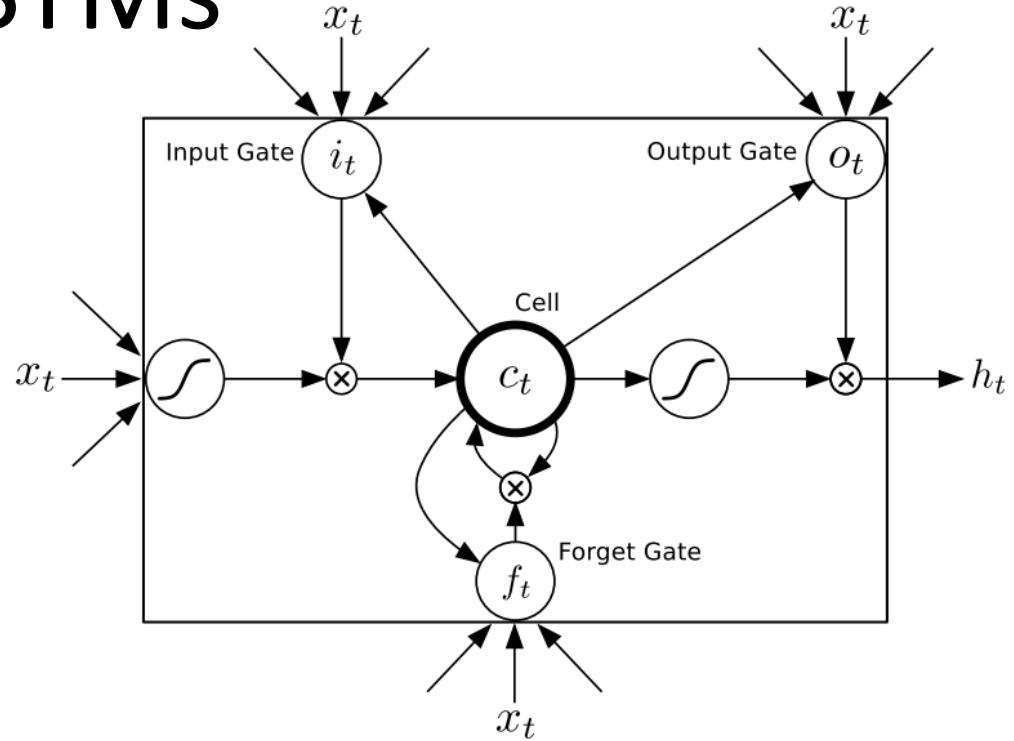
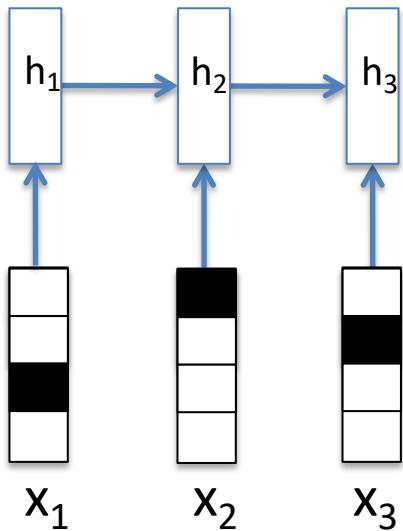
Recurrent Neural Network

- Replace

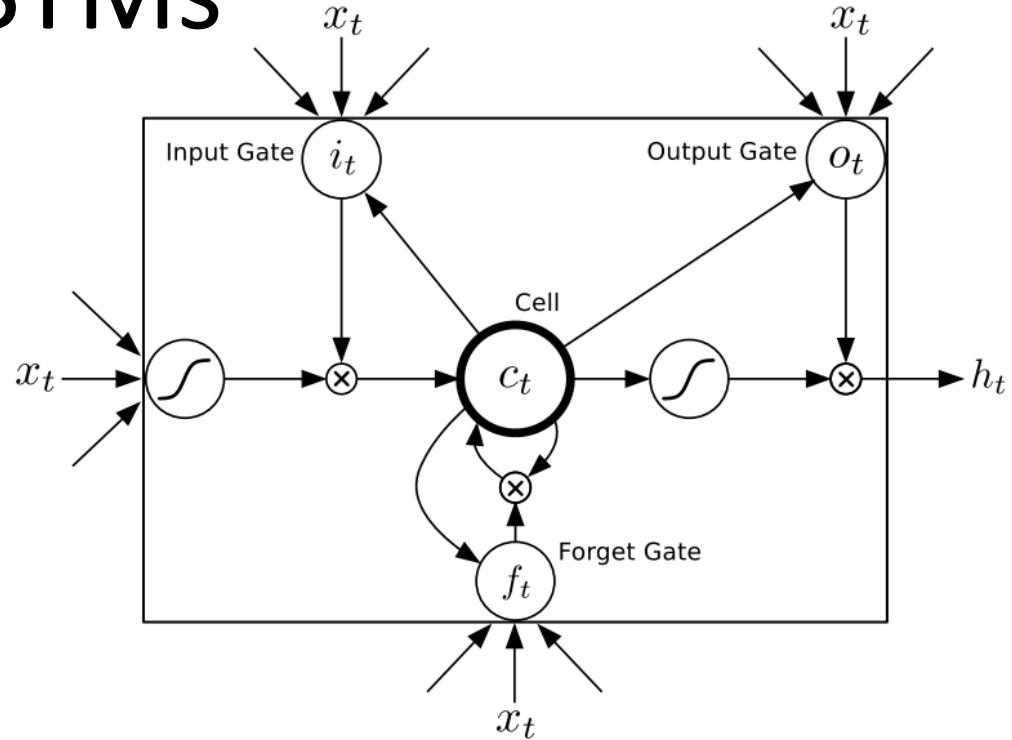
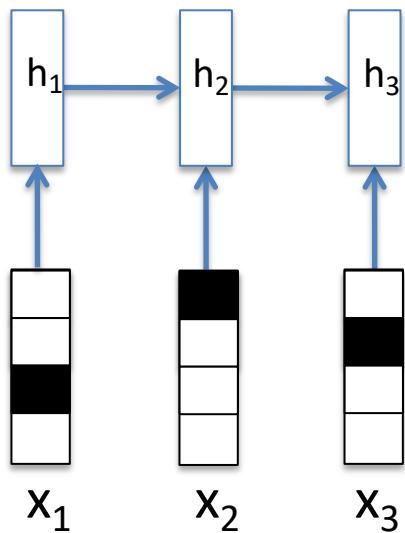


- Can be viewed as a deep neural network with tied weights.

LSTMs

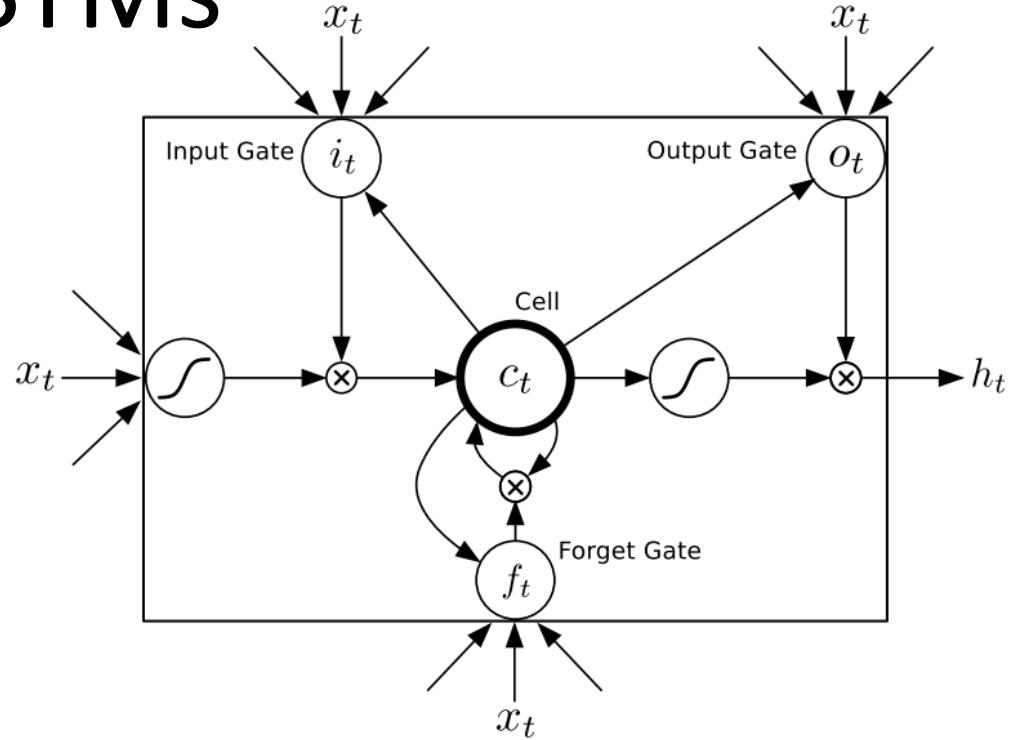
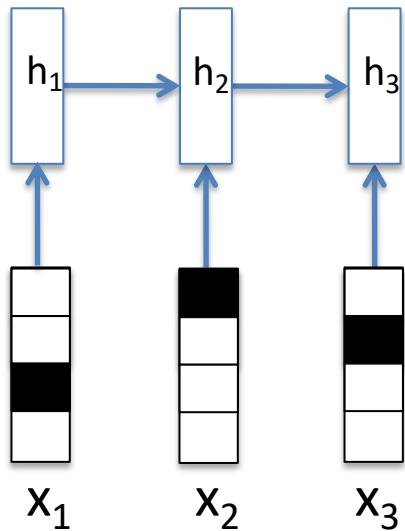


LSTMs



$$i_t = \sigma (W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i),$$

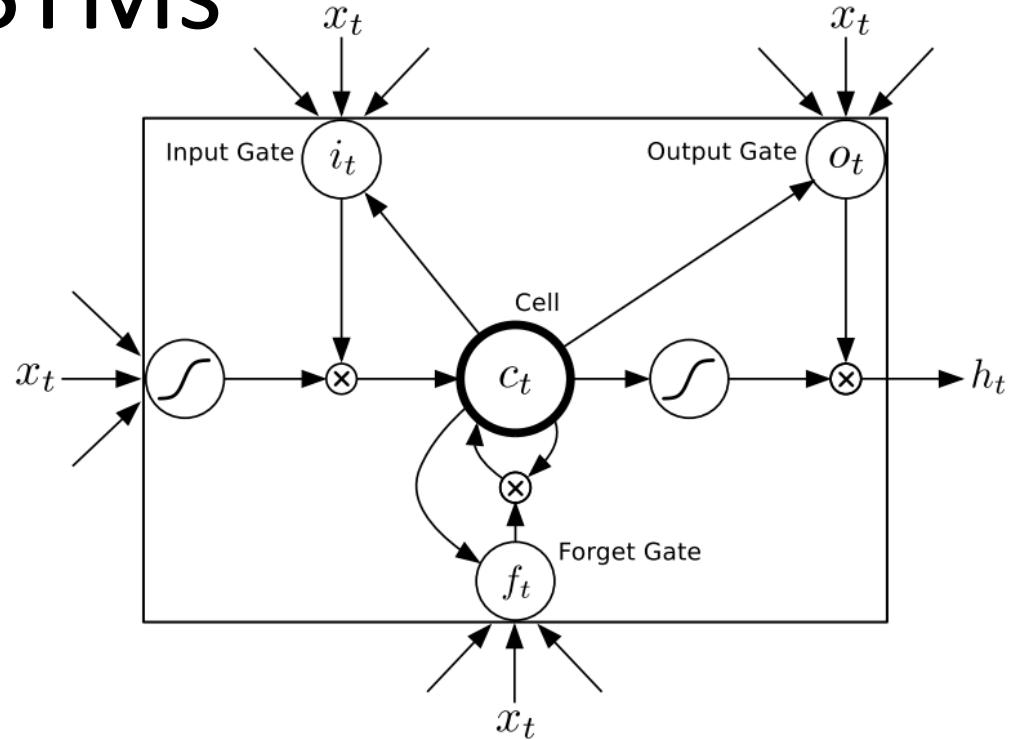
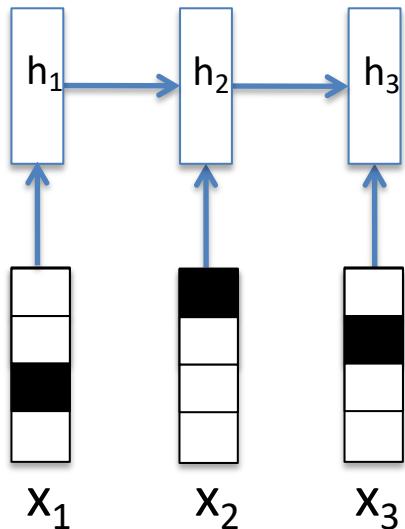
LSTMs



$$i_t = \sigma (W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i),$$

$$f_t = \sigma (W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f),$$

LSTMs

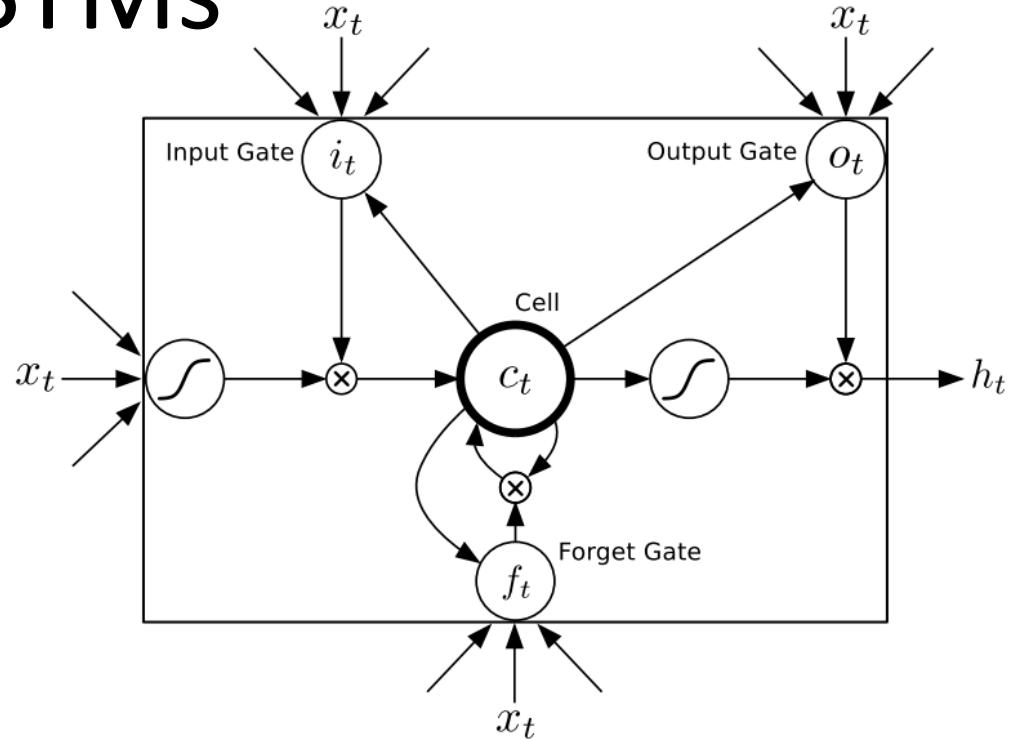
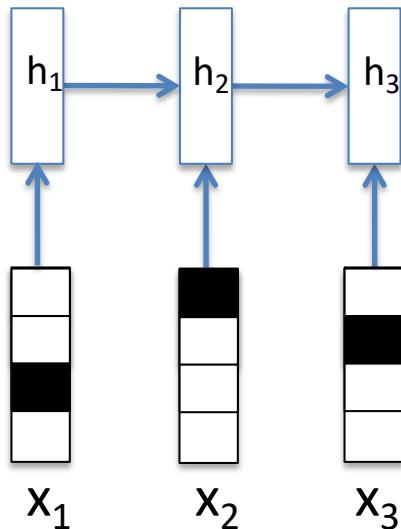


$$\mathbf{i}_t = \sigma (W_{xi} \mathbf{x}_t + W_{hi} \mathbf{h}_{t-1} + W_{ci} \mathbf{c}_{t-1} + \mathbf{b}_i),$$

$$\mathbf{f}_t = \sigma (W_{xf} \mathbf{x}_t + W_{hf} \mathbf{h}_{t-1} + W_{cf} \mathbf{c}_{t-1} + \mathbf{b}_f),$$

$$\mathbf{c}_t = \mathbf{f}_t \mathbf{c}_{t-1} + \mathbf{i}_t \tanh (W_{xc} \mathbf{x}_t + W_{hc} \mathbf{h}_{t-1} + \mathbf{b}_c),$$

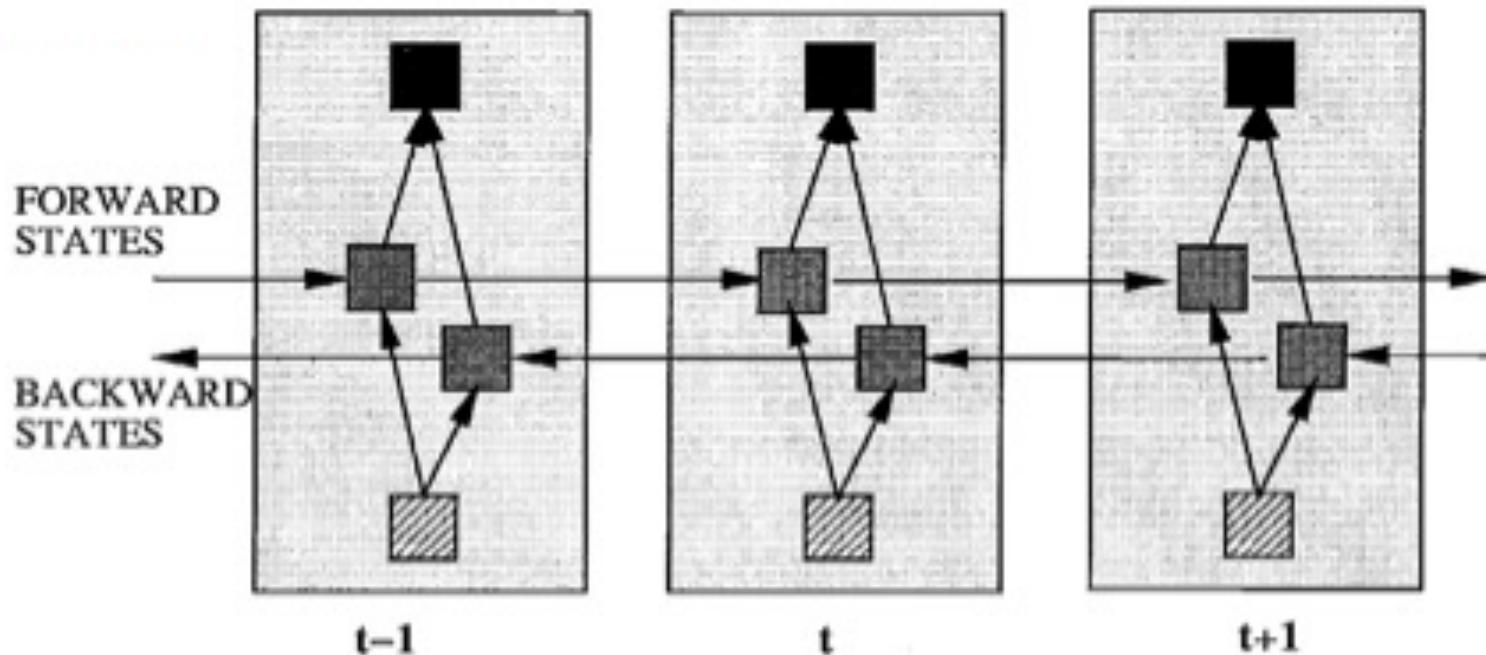
LSTMs



$$\begin{aligned}
 i_t &= \sigma(W_{xi}\mathbf{x}_t + W_{hi}\mathbf{h}_{t-1} + W_{ci}\mathbf{c}_{t-1} + \mathbf{b}_i), \\
 f_t &= \sigma(W_{xf}\mathbf{x}_t + W_{hf}\mathbf{h}_{t-1} + W_{cf}\mathbf{c}_{t-1} + \mathbf{b}_f), \\
 \mathbf{c}_t &= f_t \mathbf{c}_{t-1} + i_t \tanh(W_{xc}\mathbf{x}_t + W_{hc}\mathbf{h}_{t-1} + \mathbf{b}_c), \\
 o_t &= \sigma(W_{xo}\mathbf{x}_t + W_{ho}\mathbf{h}_{t-1} + W_{co}\mathbf{c}_t + \mathbf{b}_o), \\
 \mathbf{h}_t &= o_t \tanh(\mathbf{c}_t).
 \end{aligned}$$

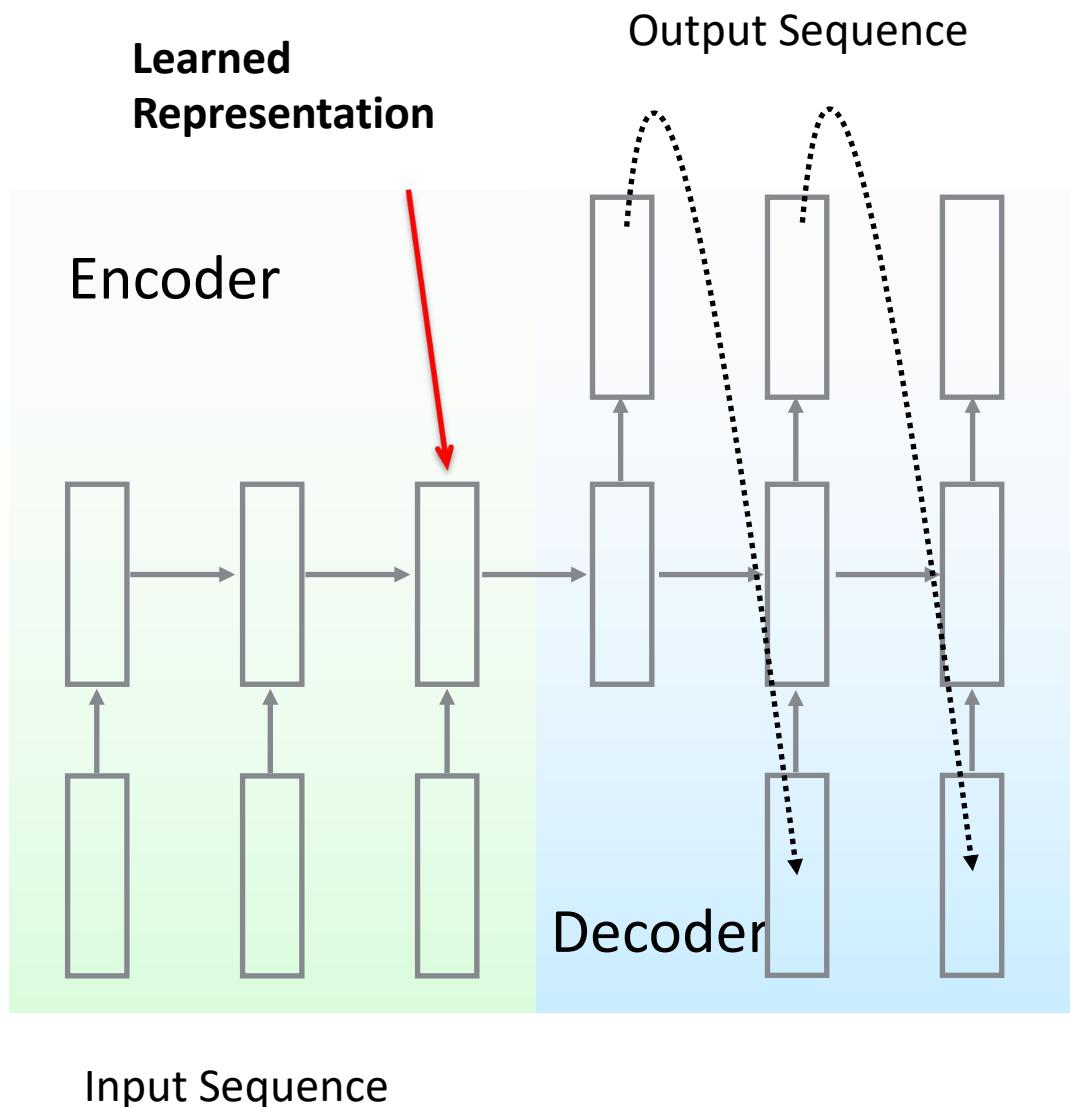
Bidirectional RNNs

Bidirectional RNNs (Schuster and Paliwal, 1997)



- Heavily used in language modeling.

Sequence to Sequence Learning

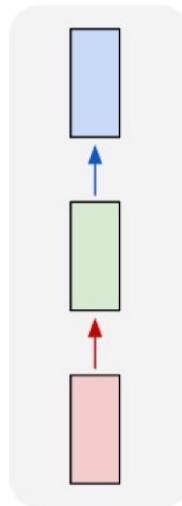


- RNN Encoder-Decoders for Machine Translation (Sutskever et al. 2014; Cho et al. 2014; Kalchbrenner et al. 2013, Srivastava et.al., 2015)

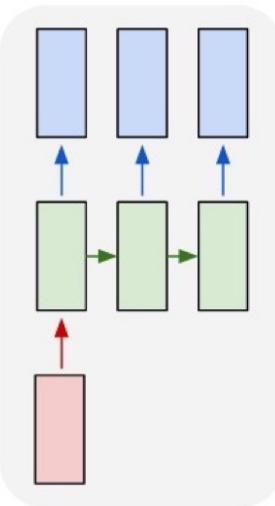
Sequence to Sequence Models

- Natural language processing is concerned with tasks involving language data

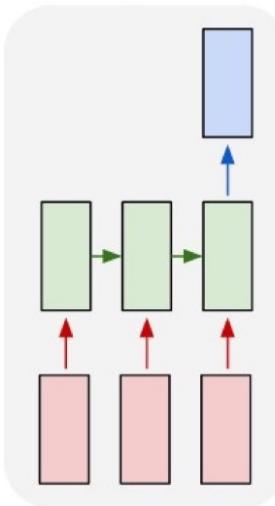
one to one



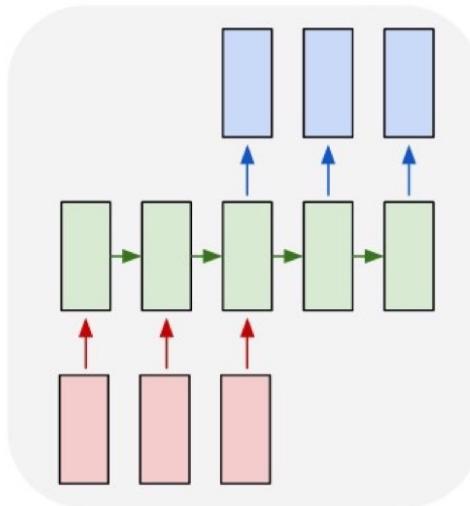
one to many



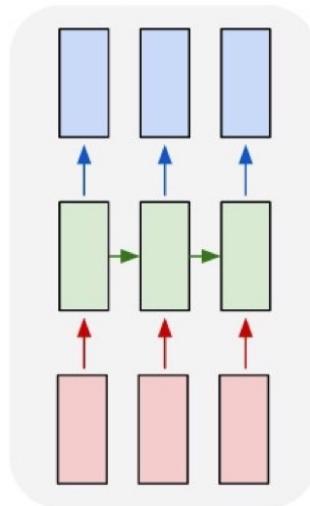
many to one



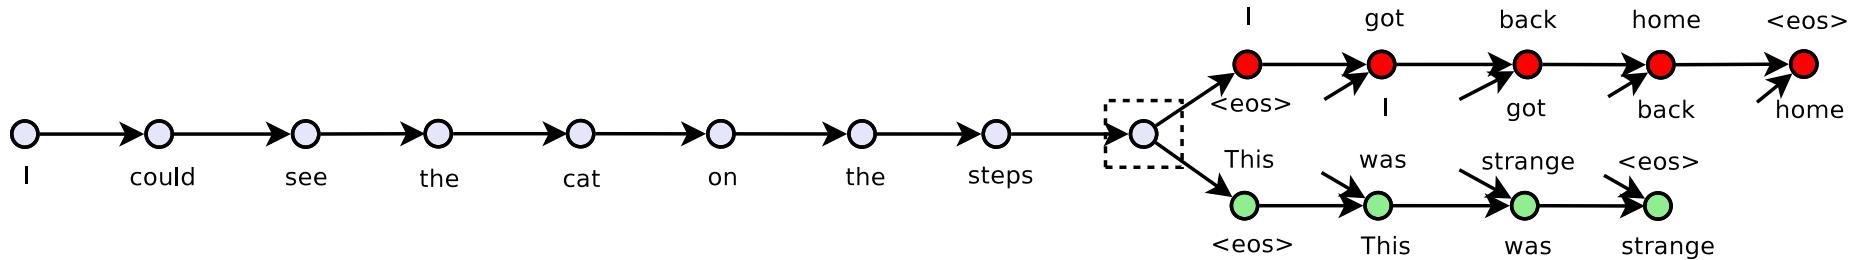
many to many



many to many

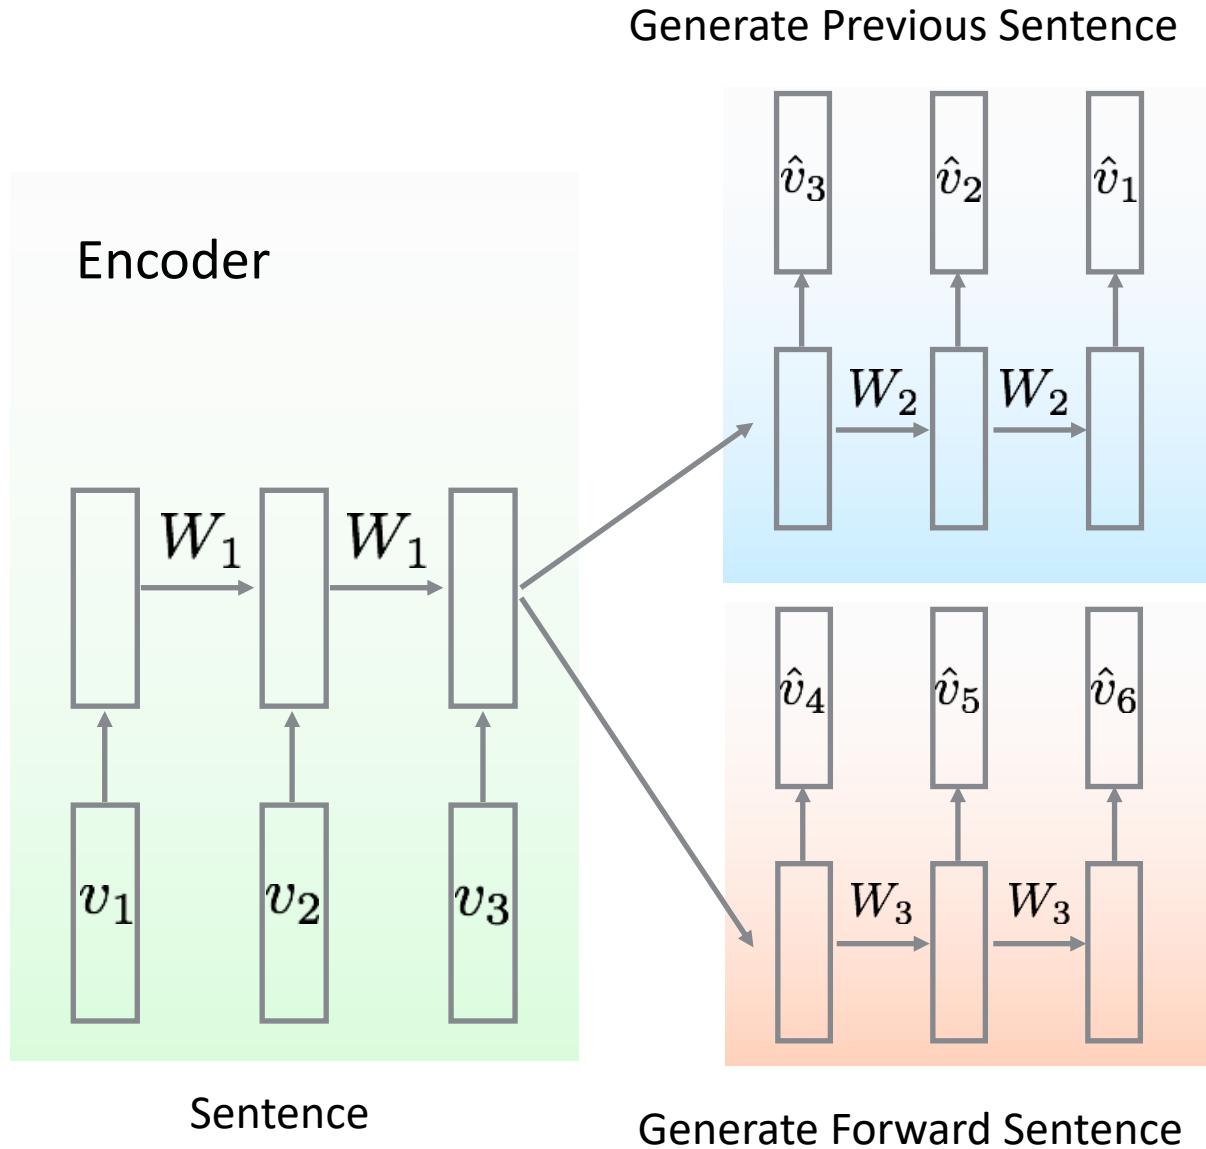


Skip-Thought Model



- Given a tuple (s_{i-1}, s_i, s_{i+1}) of contiguous sentences:
 - the sentence s_i is encoded using LSTM.
 - the sentence s_i attempts to reconstruct the previous sentence and next sentence s_{i+1} .
- The input is the sentence triplet:
 - I got back home.
 - I could see the cat on the steps.
 - This was strange.

Skip-Thought Model

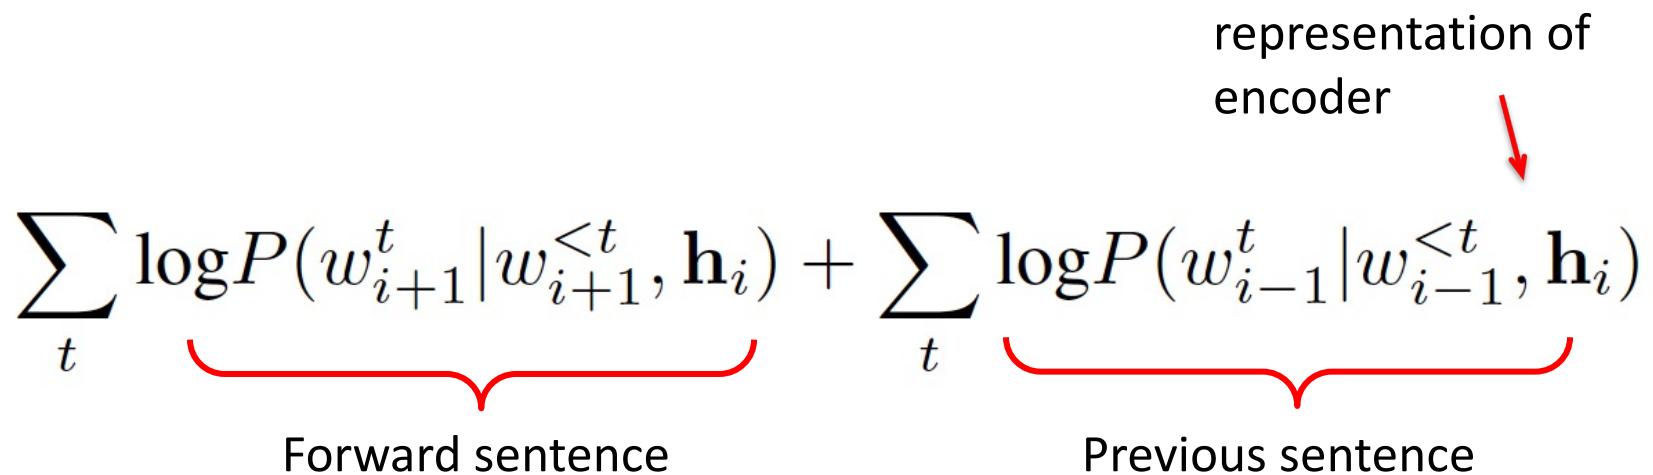


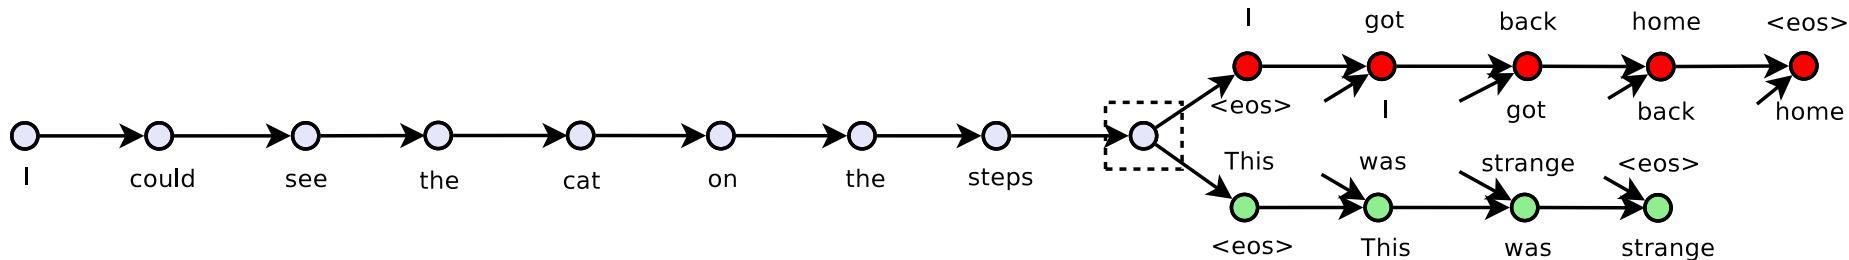
Learning Objective

- We are given a tuple (s_{i-1}, s_i, s_{i+1}) of contiguous sentences.
- **Objective:** The sum of the log-probabilities for the next and previous sentences conditioned on the encoder representation:

$$\sum_t \underbrace{\log P(w_{i+1}^t | w_{i+1}^{<t}, \mathbf{h}_i)}_{\text{Forward sentence}} + \sum_t \underbrace{\log P(w_{i-1}^t | w_{i-1}^{<t}, \mathbf{h}_i)}_{\text{Previous sentence}}$$

representation of
encoder





Book 11K corpus

# of books	# of sentences	# of words	# of unique words
11,038	74,004,228	984,846,357	1,316,420

- Query sentence along with its nearest neighbor from 500K sentences using cosine similarity:
 - He ran his hand inside his coat, double-checking that the unopened letter was still there.
 - He slipped his hand between his coat and his shirt, where the folded copies lay in a brown envelope.

Semantic Relatedness

- SemEval 2014 Task 1: semantic relatedness SICK dataset:
Given two sentences, produce a score of how semantically related these sentences are based on human generated scores (1 to 5).
- The dataset comes with a predefined split of 4500 training pairs, 500 development pairs and 4927 testing pairs.
- Using skip-thought vectors for each sentence, we simply train a linear regression to predict semantic relatedness.
 - For pair of sentences, we compute component-wise features between pairs (e.g. $|u-v|$).

Semantic Relatedness

	Method	r	ρ	MSE
SemEval 2014 sub- missions	Illinois-LH [18]	0.7993	0.7538	0.3692
	UNAL-NLP [19]	0.8070	0.7489	0.3550
	Meaning Factory [20]	0.8268	0.7721	0.3224
	ECNU [21]	0.8414	—	—
Results reported by Tai et.al.	Mean vectors [22]	0.7577	0.6738	0.4557
	DT-RNN [23]	0.7923	0.7319	0.3822
	SDT-RNN [23]	0.7900	0.7304	0.3848
	LSTM [22]	0.8528	0.7911	0.2831
	Bidirectional LSTM [22]	0.8567	0.7966	0.2736
	Dependency Tree-LSTM [22]	0.8676	0.8083	0.2532
Ours	uni-skip	0.8477	0.7780	0.2872
	bi-skip	0.8405	0.7696	0.2995
	combine-skip	0.8584	0.7916	0.2687
	combine-skip+COCO	0.8655	0.7995	0.2561

- Our models outperform all previous systems from the SemEval 2014 competition. This is remarkable, given the simplicity of our approach and the lack of feature engineering.

Semantic Relatedness

Sentence 1	Sentence 2	GT	pred
A little girl is looking at a woman in costume	A young girl is looking at a woman in costume	4.7	4.5
A little girl is looking at a woman in costume	The little girl is looking at a man in costume	3.8	4.0
A little girl is looking at a woman in costume	A little girl in costume looks like a woman	2.9	3.5
A sea turtle is hunting for fish	A sea turtle is hunting for food	4.5	4.5
A sea turtle is not hunting for fish	A sea turtle is hunting for fish	3.4	3.8
A man is driving a car	The car is being driven by a man	5	4.9
There is no man driving the car	A man is driving a car	3.6	3.5
A large duck is flying over a rocky stream	A duck, which is large, is flying over a rocky stream	4.8	4.9
A large duck is flying over a rocky stream	A large stream is full of rocks, ducks and flies	2.7	3.1
A person is performing acrobatics on a motorcycle	A person is performing tricks on a motorcycle	4.3	4.4
A person is performing tricks on a motorcycle	The performer is tricking a person on a motorcycle	2.6	4.4
Someone is pouring ingredients into a pot	Someone is adding ingredients to a pot	4.4	4.0
Nobody is pouring ingredients into a pot	Someone is pouring ingredients into a pot	3.5	4.2
Someone is pouring ingredients into a pot	A man is removing vegetables from a pot	2.4	3.6

- Example predictions from the SICK test set. GT is the ground truth relatedness, scored between 1 and 5.
- The last few results: slight changes in sentences result in large changes in relatedness that we are unable to score correctly.

Paraphrase Detection

- Microsoft Research Paraphrase Corpus: For two sentences one must predict whether or not they are paraphrases.

	Method	Acc	F1	
Recursive Auto- encoders	feats [24]	73.2		The training set contains 4076 sentence pairs (2753 are positive)
	RAE+DP [24]	72.6		
	RAE+feats [24]	74.2		
	RAE+DP+feats [24]	76.8	83.6	
Best published results	FHS [25]	75.0	82.7	The test set contains 1725 pairs (1147 are positive).
	PE [26]	76.1	82.7	
	WDDP [27]	75.6	83.0	
	MTMETRICS [28]	77.4	84.1	
Ours	uni-skip	73.0	81.9	
	bi-skip	71.2	81.2	
	combine-skip	73.0	82.0	
	combine-skip + feats	75.8	83.0	

Classification Benchmarks

- 5 datasets: movie review sentiment (MR), customer product reviews (CR), subjectivity/objectivity classification (SUBJ), opinion polarity (MPQA) and question-type classification (TREC).

	Method	MR	CR	SUBJ	MPQA	TREC
Bag-of-words	NB-SVM [41]	79.4	<u>81.8</u>	93.2	86.3	
	MNB [41]	79.0	<u>80.0</u>	<u>93.6</u>	86.3	
	cBoW [6]	77.2	79.9	<u>91.3</u>	86.4	87.3
Supervised	GrConv [6]	76.3	81.3	89.5	84.5	88.4
	RNN [6]	77.2	82.3	93.7	90.1	90.2
	BRNN [6]	82.3	82.6	94.2	90.3	91.0
	CNN [4]	81.5	85.0	93.4	89.6	93.6
	AdaSent [6]	83.1	86.3	95.5	93.3	92.4
Paragraph-vector [7]		74.8	78.1	90.5	74.2	91.8
Ours	uni-skip	75.5	79.3	92.1	86.9	91.4
	bi-skip	73.9	77.9	92.5	83.3	89.4
	combine-skip	76.5	80.1	<u>93.6</u>	87.1	<u>92.2</u>
	combine-skip + NB	<u>80.4</u>	81.3	<u>93.6</u>	<u>87.5</u>	