

10-707: Recitation on Probability Distributions

Slides adapted from Ruslan Salakhutdinov's 10-707 lecture
on probability distributions

Why do we have this recitation?

- The distributions we talk about today will be the basics of homeworks
- When we talk about generative models (GAN, VAE, Diffusion models)
 - We will generate data distributions
 - By first sampling from a probability distribution we are familiar with

Basic intro stuff

A random variable = outcome of some uncertain or random trial

RV can be discrete:

Rolling a dice $\rightarrow \{1, 2, 3, 4, 5, 6\}$ are the only possible outcomes


RV can be continuous (any real number in some interval):

Height of a person $\rightarrow 170.6475$ cm, 145.874 cm, etc.


Basic intro stuff

We have tools to describe the probability that a RV will take on a particular value

Discrete: Probability Mass Function



Continuous: Probability Density Function



Bernoulli and Binomial Distributions

Bernoulli Distribution

A Bernoulli trial models the simplest random experiment that can have **two outcomes**: success and failure

When would we use it?

- Tossing a coin and getting a heads
- Passing an exam
- Customer clicking on an ad

Bernoulli Distribution

Probability Mass Function: p is probability of 'success'

$$P(X = x) = p^x(1 - p)^{1-x}, \quad x \in \{0, 1\}$$

$$\mathbb{E}[X] = p \qquad \text{Var}(X) = p(1 - p)$$

Bernoulli Distribution

What if we don't have the value of p beforehand?

Create data from an experiment -> use it to estimate p using log likelihood!

Bernoulli Distribution – Parameter Estimation

Collect data:

$$x_1, x_2, \dots, x_n, \text{ where each } x_i \in \{0, 1\}.$$

Likelihood function:

$$L(p) = \prod_{i=1}^n p^{x_i} (1-p)^{1-x_i}$$

Log-likelihood:

$$\ell(p) = \sum_{i=1}^n [x_i \ln(p) + (1 - x_i) \ln(1 - p)]$$

Take derivative, set to 0:

$$\hat{p}_{\text{MLE}} = \frac{1}{n} \sum_{i=1}^n x_i$$

Binomial Distribution

Bernoulli: models a **single** experiment

Binomial: models a series of **n independent** trials with the **same probability of success**

When would we use it?

- Tossing a coin and getting a heads in 10 trials
- Number of customers clicking on an ad with 1000 impressions

Binomial Distribution

Probability Mass Function: p is probability of 'success'

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}, \quad k = 0, 1, 2, \dots, n$$

$$\mathbb{E}[X] = np \quad \text{Var}(X) = np(1 - p) \quad \hat{p} = \frac{\sum_j X_j}{mn}$$

MLE

Multinomial Distribution

Binomial: 2 outcomes per trial

Multinomial: generalizes, K outcomes per trial

$$\mu = (\mu_1, \mu_2, \dots, \mu_K), \quad \sum_{k=1}^K \mu_k = 1 \quad p(x \mid \mu) = \prod_{k=1}^K \mu_k^{x_k}$$

Multinomial Distribution

One-hot encoding

i.e. K= 4, correct category = 3

$$x = (0, 0, 1, 0)$$

Beta Distributions

Beta Distribution

- Bernoulli and Binomial **model data** (the probability of a random variable taking on a value)
- Beta models uncertainty about p (**the probability of p taking on a certain value**)

Beta Distribution

Critical Idea:


We have some data that we collected after running an experiment but we don't know what the probability p (success) is. So, we use the beta distribution to find the probability that p takes on a certain value between 0 and 1

Beta Distribution



7 heads, 3 tails

Coin 0	Coin 1	Coin 2	Coin 3	Coin 4	Coin 5	Coin 6	Coin 7	Coin 8	Coin 9	Coin 10
$P(H) = 0$	$P(H) = 0.1$	$P(H) = 0.2$	$P(H) = 0.3$	$P(H) = 0.4$	$P(H) = 0.5$	$P(H) = 0.6$	$P(H) = 0.7$	$P(H) = 0.8$	$P(H) = 0.9$	$P(H) = 1$
$0^7 \cdot 1^3$	$0.1^7 \cdot 0.9^3$	$0.2^7 \cdot 0.8^3$	$0.3^7 \cdot 0.7^3$	$0.4^7 \cdot 0.6^3$	$0.5^7 \cdot 0.5^3$	$0.6^7 \cdot 0.4^3$	$0.7^7 \cdot 0.3^3$	$0.8^7 \cdot 0.2^3$	$0.9^7 \cdot 0.1^3$	$1^7 \cdot 0^3$
0	0.00001	0.001	0.01	0.04	0.13	0.24	0.29	0.21	0.06	0




Beta Distribution

Beta distribution
Beta($a+1, b+1$)



$$f(x) = \frac{x^a(1-x)^b}{B(a+1,b+1)}$$



Beta Distribution

If $X \sim \text{Beta}(\alpha, \beta)$:

$$f(x; \alpha, \beta) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1}, \quad 0 < x < 1$$

where

$$B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha + \beta)}$$

$$\mathbb{E}[X] = \frac{\alpha}{\alpha + \beta} \quad \text{Var}(X) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)} \quad \Gamma(x) = \int_0^\infty u^{x-1} e^{-u} du$$

Dirichlet Distribution


- Consider a distribution over μ_k , subject to constraints:

$$\forall k : \mu_k \geq 0 \quad \text{and} \quad \sum_{k=1}^K \mu_k = 1$$

- The Dirichlet Distribution is defined as:

$$\text{Dir}(\boldsymbol{\mu}|\boldsymbol{\alpha}) = \frac{\Gamma(\alpha_0)}{\Gamma(\alpha_1) \cdots \Gamma(\alpha_K)} \prod_{k=1}^K \mu_k^{\alpha_k - 1}$$

$$\alpha_0 = \sum_{k=1}^K \alpha_k$$




Where $\alpha_1, \alpha_2, \dots, \alpha_k$ are the parameters of the distribution and $\Gamma(x)$ is the gamma function.


- The Dirichlet distribution is confined to a simplex as a consequence of the constraints.

Dirichlet Distribution


- Plots of the Dirichlet distribution over three variables.



$$\alpha_k = 10^{-1}$$




$$\alpha_k = 10^0$$



$$\alpha_k = 10^1$$

Gaussian Univariate Distribution

In the case of a single variable x , the Gaussian distribution takes form:



$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$

which is governed by two parameters

The Gaussian distribution satisfies:


$$\mathcal{N}(x|\mu, \sigma^2) > 0$$

$$\int_{-\infty}^{\infty} \mathcal{N}(x|\mu, \sigma^2) dx = 1$$

Multivariate Gaussian Distribution

- For a D-dimensional vector \mathbf{x} , the Gaussian distribution takes form:

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}$$



$\boldsymbol{\mu}$ is a D-dimensional mean vector.
and $|\boldsymbol{\Sigma}|$ denotes the determinant of $\boldsymbol{\Sigma}$.

- Note that the covariance matrix is a symmetric positive definite matrix.

Moments of the Gaussian Distribution

- The expectation of \mathbf{x} under the Gaussian distribution:

$$\begin{aligned}\mathbb{E}[\mathbf{x}] &= \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \int \exp \left\{ -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\} \mathbf{x} d\mathbf{x} \\ &= \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \int \exp \left\{ -\frac{1}{2}\mathbf{z}^T \Sigma^{-1} \mathbf{z} \right\} (\mathbf{z} + \boldsymbol{\mu}) d\mathbf{z}\end{aligned}$$


The term in \mathbf{z} in the factor $(\mathbf{z} + \boldsymbol{\mu})$ will vanish by symmetry (positive vs negative).

$$\mathbb{E}[\mathbf{x}] = \boldsymbol{\mu}$$

Moments of the Gaussian Distribution

- The second order moments of the Gaussian distribution:

$$\mathbb{E}[\mathbf{x}\mathbf{x}^T] = \boldsymbol{\mu}\boldsymbol{\mu}^T + \boldsymbol{\Sigma}$$

- The covariance is given by:

$$\text{cov}[\mathbf{x}] = \mathbb{E} [(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{x} - \mathbb{E}[\mathbf{x}])^T] = \boldsymbol{\Sigma}$$

$$\mathbb{E}[\mathbf{x}] = \boldsymbol{\mu}$$


- Because the parameter matrix $\boldsymbol{\Sigma}$ governs the covariance of \mathbf{x} under the Gaussian distribution, it is called the covariance matrix.
- Lookup table for 1d-Gaussian moments: [Normal distribution - Wikipedia](#)


Central Limit Theorem

- The distribution of mean of N i.i.d. random variables tends to a Gaussian distribution as N increases
 - sample mean approximate normality
- It allows us to make statistical inferences, validate performance, and use algorithms that assume normality
 - i.e., assume data come from a gaussian

$$\bar{X}_n \equiv \frac{X_1 + \cdots + X_n}{n}. \quad \sqrt{n} (\bar{X}_n - \mu) \xrightarrow{d} \mathcal{N}(0, \sigma^2).$$

Central Limit Theorem

- The distribution of mean of N i.i.d. random variables tends to a Gaussian distribution as N increases
 - sample mean approximate normality
- It allows us to make statistical inferences, validate performance, and use algorithms that assume normality
 - i.e., assume data come from a gaussian



Partitioned Gaussian Distribution

- Consider a D-dimensional Gaussian distribution: $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma})$
Let us partition \mathbf{x} into two disjoint subsets \mathbf{x}_a and \mathbf{x}_b :

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \end{pmatrix} \quad \boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{pmatrix} \quad \boldsymbol{\Sigma} = \begin{pmatrix} \boldsymbol{\Sigma}_{aa} & \boldsymbol{\Sigma}_{ab} \\ \boldsymbol{\Sigma}_{ba} & \boldsymbol{\Sigma}_{bb} \end{pmatrix}$$

- In many situations, it will be more convenient to work with the precision matrix (inverse of the covariance matrix):

$$\boldsymbol{\Lambda} \equiv \boldsymbol{\Sigma}^{-1} \quad \boldsymbol{\Lambda} = \begin{pmatrix} \boldsymbol{\Lambda}_{aa} & \boldsymbol{\Lambda}_{ab} \\ \boldsymbol{\Lambda}_{ba} & \boldsymbol{\Lambda}_{bb} \end{pmatrix}$$

- Note that $\boldsymbol{\Lambda}_{aa}$ is not given by the inverse of $\boldsymbol{\Sigma}_{aa}$

Conditional Distribution

- It turns out that the conditional distribution is also a Gaussian distribution:

$$p(\mathbf{x}_a | \mathbf{x}_b) = \mathcal{N}(\mathbf{x}_a | \boldsymbol{\mu}_{a|b}, \boldsymbol{\Sigma}_{a|b})$$

Covariance does not



depend on \mathbf{x}_b .

$$\begin{aligned}\boldsymbol{\Sigma}_{a|b} &= \boldsymbol{\Lambda}_{aa}^{-1} = \boldsymbol{\Sigma}_{aa} - \boldsymbol{\Sigma}_{ab} \boldsymbol{\Sigma}_{bb}^{-1} \boldsymbol{\Sigma}_{ba} \\ \boldsymbol{\mu}_{a|b} &= \boldsymbol{\Sigma}_{a|b} \{ \boldsymbol{\Lambda}_{aa} \boldsymbol{\mu}_a - \boldsymbol{\Lambda}_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b) \} \\ &= \boldsymbol{\mu}_a - \boldsymbol{\Lambda}_{aa}^{-1} \boldsymbol{\Lambda}_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b) \\ &= \boldsymbol{\mu}_a + \boldsymbol{\Sigma}_{ab} \boldsymbol{\Sigma}_{bb}^{-1} (\mathbf{x}_b - \boldsymbol{\mu}_b)\end{aligned}$$



Linear function
of \mathbf{x}_b .

Marginal Distribution

- The marginal distribution is also a Gaussian distribution:

$$\begin{aligned} p(\mathbf{x}_a) &= \int p(\mathbf{x}_a, \mathbf{x}_b) d\mathbf{x}_b \\ &= \mathcal{N}(\mathbf{x}_a | \boldsymbol{\mu}_a, \boldsymbol{\Sigma}_{aa}) \end{aligned}$$

- For a marginal distribution, the mean and covariance are most simply expressed in terms of partitioned covariance matrix.

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \end{pmatrix}$$

$$\boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{pmatrix}$$

$$\boldsymbol{\Sigma} = \begin{pmatrix} \boldsymbol{\Sigma}_{aa} & \boldsymbol{\Sigma}_{ab} \\ \boldsymbol{\Sigma}_{ba} & \boldsymbol{\Sigma}_{bb} \end{pmatrix}$$

Maximum Likelihood Estimation

- Suppose we observed i.i.d data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$.
- We can construct the log-likelihood function, which is a function of $\boldsymbol{\mu}$ and Σ :

$$\ln p(\mathbf{X}|\boldsymbol{\mu}, \Sigma) = -\frac{ND}{2} \ln(2\pi) - \frac{N}{2} \ln |\Sigma| - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu})^T \Sigma^{-1} (\mathbf{x}_n - \boldsymbol{\mu})$$

- Note that the likelihood function depends on the N data points only through the following sums:

$$\sum_{n=1}^N \mathbf{x}_n$$

$$\sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^T$$

Sufficient Statistics

Maximum Likelihood Estimation

- To find a maximum likelihood estimate of the mean, we set the derivative of the log-likelihood function to zero:

$$\frac{\partial}{\partial \boldsymbol{\mu}} \ln p(\mathbf{X} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sum_{n=1}^N \boldsymbol{\Sigma}^{-1} (\mathbf{x}_n - \boldsymbol{\mu}) = 0$$

and solve to obtain:

$$\boldsymbol{\mu}_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n.$$

- Similarly, we can find the maximum likelihood estimate of $\boldsymbol{\Sigma}$

$$\boldsymbol{\Sigma}_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})(\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})^T.$$

Maximum Likelihood Estimation

- Evaluating the expectation of the maximum likelihood estimates under the true distribution, we obtain:

$$\mathbb{E}[\boldsymbol{\mu}_{\text{ML}}] = \boldsymbol{\mu} \quad \xrightarrow{\text{Unbiased estimate}}$$

$$\mathbb{E}[\boldsymbol{\Sigma}_{\text{ML}}] = \frac{N-1}{N} \boldsymbol{\Sigma}. \quad \xrightarrow{\text{Biased estimate}}$$

Note that the maximum likelihood estimate of $\boldsymbol{\Sigma}$ is biased.

We can correct the bias by defining a different estimator:

$$\tilde{\boldsymbol{\Sigma}} = \frac{1}{N-1} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})(\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})^T.$$